

Wikipedia Text for Land-Cover

Key Questions & Experimental Findings

1. Text Signal vs. Shuffled

Q: Does the text actually contain land-cover signal, or is it just an effect of priors?

A: Yes. Aligned text significantly beats shuffled text, especially for CORINE.

2. Changing the LLM Classifier

Q: Does the result depend on a single LLM?

A: No. Both Qwen and Gemma extract signal. Qwen excels on balanced metrics; Gemma favors majority classes.

3. Spatial Confidence

Q: Is the geographic point-in-polygon label reliable?

A: Not always. Many points fall in mixed zones, requiring spatial purity measurement around the point.

4. Semantic Confidence

Q: Are all Wikipedia articles useful for predicting land-cover?

A: No. Articles with high landscape and land-cover relevance give much more robust results.

5. Text Transformation

Q: Can raw content be replaced by summaries, cards, highlights, or retrieved windows?

A: Not really. These transformations help with diagnostics and filtering, but raw content remains the best direct input.